

Linear Regression

Md. Aminul Islam Shazid

Outline

- 1 Data as Vectors and Matrices
- 2 Simple Linear Regression
- 3 Multiple Linear Regression
- 4 Prediction Interval
- 5 Qualitative Predictors
- 6 Interaction

Data as Vectors and Matrices

Observations, features and labels

i	features			label
	TV	radio	newspaper	sales
1	230.1	37.8	69.2	22.1
2	44.5	39.3	45.1	10.4
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$
200	232.1	8.6	8.7	13.4

- Each row is one **observation** (a market): $n = 200$ observations
- The columns TV, radio, newspaper are the **features** (predictors, inputs, independent variables): budgets in thousands of dollars
- Number of columns/features/predictors: $p = 3$
- The column sales is the **label** (response, output, target, dependent variable): what we want to predict, in thousands of units

Vectors and matrices

- Observation i is a **feature vector** $\mathbf{x}^{(i)} = (x_1^{(i)}, x_2^{(i)}, \dots, x_p^{(i)})^\top \in \mathbb{R}^p$ with label $y^{(i)} \in \mathbb{R}$; for example $\mathbf{x}^{(1)} = (230.1, 37.8, 69.2)^\top$ and $y^{(1)} = 22.1$
- Stacking the n observations gives the **feature matrix** $\mathbf{X} \in \mathbb{R}^{n \times p}$, whose row i is $(\mathbf{x}^{(i)})^\top$, and the **label vector** $\mathbf{y} \in \mathbb{R}^n$

$$\mathbf{X} = \begin{bmatrix} 230.1 & 37.8 & 69.2 \\ 44.5 & 39.3 & 45.1 \\ \vdots & \vdots & \vdots \\ 232.1 & 8.6 & 8.7 \end{bmatrix} = \begin{bmatrix} (\mathbf{x}^{(1)})^\top \\ (\mathbf{x}^{(2)})^\top \\ \vdots \\ (\mathbf{x}^{(n)})^\top \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 22.1 \\ 10.4 \\ \vdots \\ 13.4 \end{bmatrix}$$

- Superscript (i) = which observation (row), subscript j = which feature (column): $x_j^{(i)}$ is entry (i, j) of $\mathbf{X}$; bold lowercase = column vector, bold uppercase = matrix, plain letter = single number


The regression problem

- Assume the label depends on the features through $y = f(\mathbf{x}) + \varepsilon$: f is a fixed but unknown function and ε is a random error with mean zero, independent of $\mathbf{x}$
- **Supervised learning**: estimate f from the labelled data $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^n$, which gives an estimate $\hat{f}$
- **Regression** when the label is quantitative (sales, mpg, balance); classification when it is qualitative (later lectures)
- **Prediction**: compute $\hat{y} = \hat{f}(\mathbf{x})$ for new observations
- **Inference**: understand how the label changes with each feature (which features matter, how large the effects are, whether the relationship is linear)

Advertising data: questions regression can answer

- Is there a relationship between the advertising budgets and sales, and how strong is it?
- Which media are associated with sales, and by how much does each contribute?
- How accurately can we predict sales for given budgets?
- Is the relationship linear, and do the media reinforce each other (synergy)?

Advertising data: questions regression can answer (cont.)



Simple Linear Regression

Simple linear regression model

- One feature x and a quantitative label y , related by a straight line plus error: $y = \beta_0 + \beta_1 x + \varepsilon$
- β_0 is the **intercept** (the value of y when $x = 0$), β_1 is the **slope** (the change in y for a one-unit increase in x); together they are the **coefficients**
- Estimate the coefficients from the training data (hat = estimated): $\hat{\beta}_0, \hat{\beta}_1$; predict with $\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$
- **Fitted value** of observation i : $\hat{y}^{(i)} = \hat{\beta}_0 + \hat{\beta}_1 x^{(i)}$; **residual**: $e^{(i)} = y^{(i)} - \hat{y}^{(i)}$; residuals are computed from data, the errors $\varepsilon^{(i)}$ are never observed
- Example: $\text{sales} \approx \beta_0 + \beta_1 \times \text{TV}$

Assumptions of the linear regression model

- For $i = 1, \dots, n$: $y^{(i)} = \beta_0 + \beta_1 x^{(i)} + \varepsilon^{(i)}$, with the features $x^{(i)}$ treated as fixed numbers
- **Linearity, zero-mean errors:** $\mathbb{E}[\varepsilon^{(i)}] = 0$, so $\mathbb{E}[y | x] = \beta_0 + \beta_1 x$; this line is the *population regression line*
- **Constant variance, uncorrelated errors:** $\text{Var}(\varepsilon^{(i)}) = \sigma^2$ for every i and $\text{Cov}(\varepsilon^{(i)}, \varepsilon^{(k)}) = 0$ for $i \neq k$
- **Normality**, needed only for exact small-sample tests and intervals: $\varepsilon^{(i)} \sim N(0, \sigma^2)$
- Computing $\hat{\beta}_0, \hat{\beta}_1$ needs none of these; unbiasedness and the standard errors use the first three; how to check them with residuals is the topic of the next lecture



Least squares criterion

- **Residual sum of squares:**

$$\text{RSS} = \sum_{i=1}^n (e^{(i)})^2 = \sum_{i=1}^n (y^{(i)} - \beta_0 - \beta_1 x^{(i)})^2$$

- **Least squares** chooses $\hat{\beta}_0, \hat{\beta}_1$ to minimize RSS: the line with the smallest total squared vertical distance to the points
- Squaring stops positive and negative residuals from cancelling and penalizes large ones more; the red dot marks the minimum (RSS = 2102.5 for TV)

Least squares criterion (cont.)



OLS for simple regression: first-order conditions

Minimize $\text{RSS}(\beta_0, \beta_1) = \sum_{i=1}^n (y^{(i)} - \beta_0 - \beta_1 x^{(i)})^2$. At the minimum both partial derivatives are zero:

$$\frac{\partial \text{RSS}}{\partial \beta_0} = -2 \sum_{i=1}^n (y^{(i)} - \beta_0 - \beta_1 x^{(i)}) = 0$$

$$\frac{\partial \text{RSS}}{\partial \beta_1} = -2 \sum_{i=1}^n x^{(i)} (y^{(i)} - \beta_0 - \beta_1 x^{(i)}) = 0$$

Rearranging gives the **normal equations**, two linear equations in the two unknowns β_0, β_1 :

$$\begin{aligned} n \beta_0 + \left(\sum_i x^{(i)} \right) \beta_1 &= \sum_i y^{(i)} \\ \left(\sum_i x^{(i)} \right) \beta_0 + \left(\sum_i (x^{(i)})^2 \right) \beta_1 &= \sum_i x^{(i)} y^{(i)} \end{aligned}$$

OLS for simple regression: solving the normal equations

- Dividing the first normal equation by n gives $\bar{y} = \beta_0 + \beta_1 \bar{x}$, so

$$\boxed{\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}}$$

- Substituting into the second gives

$$\begin{aligned}\sum_i x^{(i)}(y^{(i)} - \bar{y}) &= \hat{\beta}_1 \sum_i x^{(i)}(x^{(i)} - \bar{x}) \\ \hat{\beta}_1 &= \frac{\sum_i x^{(i)}(y^{(i)} - \bar{y})}{\sum_i x^{(i)}(x^{(i)} - \bar{x})} \\ &= \frac{\sum_i (x^{(i)} - \bar{x})(y^{(i)} - \bar{y})}{\sum_i (x^{(i)} - \bar{x})^2} \\ &= \frac{S_{xy}}{S_{xx}}\end{aligned}$$

Interpreting the coefficients

- Fitted line for the Advertising data: $\widehat{\text{sales}} = 7.03 + 0.0475 \times \text{TV}$ (sales in thousands of units, TV budget in thousands of dollars)
- $\hat{\beta}_1 = 0.0475$: an additional \$1,000 of TV budget is associated with about $0.0475 \times 1000 = 47.5$ additional units sold
- $\hat{\beta}_0 = 7.03$: about 7,030 units are predicted when the TV budget is zero; the smallest TV budget in the data is \$700, so the intercept sits at the edge of the data
- The fit describes association in the observed markets; a causal reading needs more than a fitted line, because budgets were not assigned at random

Properties of the least squares line


- The line passes through the point of means $(\bar{x}, \bar{y})$, by the first normal equation; Advertising: $7.033 + 0.04754 \times 147.04 = 14.02 = \bar{y}$
- The normal equations say $\sum_i e^{(i)} = 0$ and $\sum_i x^{(i)} e^{(i)} = 0$: the residuals sum to zero and are orthogonal to the feature; consequently $\sum_i \hat{y}^{(i)} e^{(i)} = 0$ and the fitted values have mean $\bar{y}$
- The slope is the correlation rescaled: $\hat{\beta}_1 = r s_y / s_x$, where r is the sample correlation and s_x, s_y are the sample standard deviations; Advertising: $0.7822 \times 5.217 / 85.85 = 0.0475$
- Since $|r| \leq 1$, a one-standard-deviation increase in x goes with an increase of only r standard deviations in $\hat{y}$: the origin of the name *regression* (towards the mean)

How accurate are the estimates?


- The population regression line (red) is the true, unknown relationship; the least squares line (blue) is computed from one sample, and every new sample gives a different line
- On average the estimates are correct (unbiased), but a single $\hat{\beta}_1$ differs from β_1 ; the **standard error** measures the typical size of this difference

How accurate are the estimates? (cont.)

one sample: true line (red), fitted line (blue)



ten samples: ten different fitted lines



Mean and variance of the slope estimate

- **Unbiased:** $\mathbb{E}[\hat{\beta}_1] = \beta_1$
- **Variance:** $\text{Var}(\hat{\beta}_1) = \sigma^2 / S_{xx}$

Variance of the estimators

$$V(\hat{\beta}_1) = \frac{\sigma^2}{\sum_{i=1}^n (x^{(i)} - \bar{x})^2}, \quad V(\hat{\beta}_0) = \sigma^2 \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x^{(i)} - \bar{x})^2} \right]$$

- Here $\sigma^2 = \text{Var}(\varepsilon)$ is the variance of the errors
- σ^2 is unknown and is estimated from the residuals
- These formulas assume the errors are uncorrelated with constant variance σ^2
- The standard errors (SE) of the estimates are found taking the square roots of the variances

Estimating the error variance

- The errors $\varepsilon^{(i)}$ are unobserved, so σ^2 is estimated from the residuals $e^{(i)}$
- An unbiased estimator of σ^2 is given by the Mean Squared Error,
$$\text{MSE} = \boxed{\hat{\sigma}^2 = \text{RSS}/(n - 2)}$$
- Recall, Residual Sum of Squares,
$$\text{RSS} = \sum_{i=1}^n (e^{(i)})^2 = \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2$$
- The residual standard error (RSE) is then given by
$$\text{RSE} = \sqrt{\text{MSE}} = \sqrt{\text{RSS}/(n - 2)}$$
- Two parameters were estimated from the data, so two degrees of freedom are used up; with p features the divisor becomes $n - p - 1$
- Advertising, sales on TV: $\text{RSS} = 2102.53$ and $n - 2 = 198$ give $\hat{\sigma}^2 = 10.62$ and $\text{RSE} = 3.26$

Standard Error of the estimators

$$\text{SE}(\hat{\beta}_1) = \sqrt{\frac{\text{MSE}}{\sum_{i=1}^n (x^{(i)} - \bar{x})^2}},$$

$$\text{SE}(\hat{\beta}_0) = \sqrt{\text{MSE} \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x^{(i)} - \bar{x})^2} \right]}$$

Confidence intervals and hypothesis tests (recap)

- **95% confidence interval** for β_1 : approximately $\hat{\beta}_1 \pm t_{n-2, \frac{\alpha}{2}} \cdot \text{SE}(\hat{\beta}_1)$
- **Hypothesis test**: $H_0: \beta_1 = 0$ (no relationship) against $H_a: \beta_1 \neq 0$, using $t = \hat{\beta}_1 / \text{SE}(\hat{\beta}_1)$, which follows a t_{n-2} distribution if H_0 is true
- The **p -value** is the probability of a $|t|$ at least this large if H_0 were true; a small p -value (say below 0.05) is evidence against H_0
- Advertising, sales on TV: 95% confidence interval for β_1 is [0.0422, 0.0528]

	coefficient	std. error	t -statistic	p -value
intercept	7.0326	0.4578	15.36	< 0.0001
TV	0.0475	0.0027	17.67	< 0.0001

Assessing the fit: RSE and R-squared

- **Residual standard error:** $RSE = \sqrt{RSS/(n-2)}$, the typical size of a residual, in the units of y ; Advertising: $RSE = 3.26$, about 3,260 units (23% of the mean sales)
- R^2 : the proportion of the variation in y explained by the regression

$$R^2 = 1 - \frac{RSS}{TSS}, \quad TSS = \sum_{i=1}^n (y^{(i)} - \bar{y})^2$$

- $0 \leq R^2 \leq 1$; Advertising: $R^2 = 1 - 2102.5/5417.2 = 0.612$, so TV explains about 61% of the variance in sales
- In simple regression $R^2 = r^2$, the squared correlation between x and y ($r = 0.782$)
- RSE is in the units of y , R^2 is unit-free

Analysis of variance: decomposing the total variation

- Split each deviation from the mean into an explained and an unexplained part: $y^{(i)} - \bar{y} = (\hat{y}^{(i)} - \bar{y}) + (y^{(i)} - \hat{y}^{(i)})$
- Square and sum over i ; the cross term $2 \sum_i (\hat{y}^{(i)} - \bar{y}) e^{(i)}$ is zero because $\sum_i e^{(i)} = 0$ and $\sum_i \hat{y}^{(i)} e^{(i)} = 0$ (normal equations), which leaves

$$\underbrace{\sum_{i=1}^n (y^{(i)} - \bar{y})^2}_{\text{TSS}} = \underbrace{\sum_{i=1}^n (\hat{y}^{(i)} - \bar{y})^2}_{\text{ESS}} + \underbrace{\sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2}_{\text{RSS}}$$

- **TSS**: total sum of squares (total variation in y); **ESS**: explained (regression) sum of squares; **RSS**: residual sum of squares
- Hence $R^2 = \text{ESS}/\text{TSS} = 1 - \text{RSS}/\text{TSS}$; Advertising, sales on TV:
 $5417.15 = 3314.62 + 2102.53$

The ANOVA table for simple regression

source	df	sum of squares	mean square	F	p -value
regression (TV)	1	3314.62	3314.62	312.14	< 0.0001
residual	198	2102.53	10.62		
total	199	5417.15			

- Degrees of freedom add up: total $n - 1 =$ regression 1 + residual $n - 2$; a mean square is a sum of squares divided by its df, so $MS_{\text{res}} = RSE^2 = 10.62$
- $F = MS_{\text{reg}}/MS_{\text{res}}$ tests $H_0: \beta_1 = 0$; with normal errors $F \sim F_{1, n-2}$ under H_0

Multiple Linear Regression

Why not several simple regressions?

sales on	$\hat{\beta}_1$	std. error	t -statistic	p -value
TV	0.0475	0.0027	17.67	< 0.0001
radio	0.2025	0.0204	9.92	< 0.0001
newspaper	0.0547	0.0166	3.30	0.0011

- Three separate simple regressions give no single prediction from all three budgets, and each one ignores the other media
- A simple-regression slope also absorbs the effect of every correlated feature that was left out
- Newspaper and radio budgets are correlated ($r = 0.354$): markets that spend more on radio tend to spend more on newspaper, so newspaper may look useful only because it stands in for radio
- Remedy: put all features in one model

Multiple linear regression model

- With p features: $y = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p + \varepsilon$; β_j is the average change in y for a one-unit increase in x_j , *holding all other features fixed*
- For all n observations at once, $\mathbf{y} = \tilde{\mathbf{X}}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$:

$$\underbrace{\begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(n)} \end{bmatrix}}_{\mathbf{y}} = \underbrace{\begin{bmatrix} 1 & x_1^{(1)} & \dots & x_p^{(1)} \\ \vdots & \vdots & & \vdots \\ 1 & x_1^{(n)} & \dots & x_p^{(n)} \end{bmatrix}}_{\tilde{\mathbf{X}}} \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{bmatrix}}_{\boldsymbol{\beta}} + \underbrace{\begin{bmatrix} \varepsilon^{(1)} \\ \vdots \\ \varepsilon^{(n)} \end{bmatrix}}_{\boldsymbol{\varepsilon}}$$

- $\tilde{\mathbf{X}} = [\mathbf{1} \ \mathbf{X}] \in \mathbb{R}^{n \times (p+1)}$ is the feature matrix with a column of ones for the intercept (the **design matrix**); fitted values $\hat{\mathbf{y}} = \tilde{\mathbf{X}}\hat{\boldsymbol{\beta}}$, residuals $\mathbf{e} = \mathbf{y} - \hat{\mathbf{y}}$

OLS for multiple regression: normal equations

Minimize $\text{RSS}(\beta) = \|\mathbf{y} - \tilde{\mathbf{X}}\beta\|^2 = (\mathbf{y} - \tilde{\mathbf{X}}\beta)^\top (\mathbf{y} - \tilde{\mathbf{X}}\beta)$. Expanding, and using that the scalar $\beta^\top \tilde{\mathbf{X}}^\top \mathbf{y}$ equals its transpose $\mathbf{y}^\top \tilde{\mathbf{X}}\beta$:

$$\text{RSS}(\beta) = \mathbf{y}^\top \mathbf{y} - 2\beta^\top \tilde{\mathbf{X}}^\top \mathbf{y} + \beta^\top \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \beta$$

With the rules $\nabla_\beta(\mathbf{a}^\top \beta) = \mathbf{a}$ and $\nabla_\beta(\beta^\top A\beta) = 2A\beta$ for symmetric A , the gradient is

$$\nabla_\beta \text{RSS} = -2\tilde{\mathbf{X}}^\top \mathbf{y} + 2\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \beta = \mathbf{0} \quad \implies \quad \tilde{\mathbf{X}}^\top \tilde{\mathbf{X}} \hat{\beta} = \tilde{\mathbf{X}}^\top \mathbf{y}$$

These are the **normal equations**: $p + 1$ linear equations in the $p + 1$ unknowns $\beta_0, \dots, \beta_p$.

OLS for multiple regression: solution

- If $\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}}$ is invertible, the normal equations have the unique solution

$$\hat{\boldsymbol{\beta}} = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \mathbf{y}$$

Sampling distribution of the OLS estimator

- Substitute $\mathbf{y} = \tilde{\mathbf{X}}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$ into the solution: $\hat{\boldsymbol{\beta}} = \boldsymbol{\beta} + (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top \boldsymbol{\varepsilon}$, so $\hat{\boldsymbol{\beta}}$ is the truth plus a linear function of the errors
- **Unbiased:** $\mathbb{E}[\boldsymbol{\varepsilon}] = \mathbf{0}$ gives $\mathbb{E}[\hat{\boldsymbol{\beta}}] = \boldsymbol{\beta}$
- **Covariance:** with $\text{Cov}(\boldsymbol{\varepsilon}) = \sigma^2 \mathbf{I}$,

$$\text{Cov}(\hat{\boldsymbol{\beta}}) = (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^\top (\sigma^2 \mathbf{I}) \tilde{\mathbf{X}} (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} = \sigma^2 (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1}$$

- If errors follow the normal distribution, then $\hat{\boldsymbol{\beta}} \sim N(\boldsymbol{\beta}, \sigma^2 (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1})$;
 $\text{SE}(\hat{\beta}_j) = \hat{\sigma} \sqrt{[(\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1}]_{jj}}$ with $\hat{\sigma}^2 = \text{RSS}/(n - p - 1)$, and
 $t_j = \hat{\beta}_j / \text{SE}(\hat{\beta}_j) \sim t_{n-p-1}$ under $H_0: \beta_j = 0$

Interpreting the coefficients

	coefficient	std. error	<i>t</i> - statistic	<i>p</i> - value
intercept	2.9389	0.3119	9.42	< 0.0001
TV	0.0458	0.0014	32.81	< 0.0001
radio	0.1885	0.0086	21.89	< 0.0001
newspaper	-0.0010	0.0059	-0.18	0.860

- Holding radio and newspaper fixed, an extra \$1,000 on TV is associated with about 45.8 more units sold; an extra \$1,000 on radio with about 188.5 more units
- Newspaper is not significant ($p = 0.86$) once TV and radio are in the model, although it was significant on its own ($p = 0.0011$)
- The correlation between newspaper and radio explains it: the simple regression credited newspaper with the effect of radio

Is at least one feature useful? The F-test

- $H_0: \beta_1 = \beta_2 = \dots = \beta_p = 0$ against H_a : at least one $\beta_j \neq 0$

$$F = \frac{(\text{TSS} - \text{RSS})/p}{\text{RSS}/(n - p - 1)} = \frac{\text{ESS}/p}{\text{RSS}/(n - p - 1)}$$

- If H_0 is true, F is close to 1; a value much larger than 1 is evidence against H_0 , with the p -value taken from the $F_{p, n-p-1}$ distribution
- Advertising: $F = 570.3$ on $(3, 196)$ degrees of freedom, $p < 0.0001$: at least one medium is related to sales
- Why not just look at the individual t -tests? With many features, about 5% of the p -values fall below 0.05 by chance even if no feature matters; the F -test accounts for the number of features
- Deciding *which* features to keep is variable selection, deferred to a later lecture

The ANOVA table for multiple regression

source	df	sum of squares	mean square	F	p -value
regression	3	4860.32	1620.11	570.27	< 0.0001
residual	196	556.83	2.84		
total	199	5417.15			

- Same layout as in simple regression, now with all three media: regression has p degrees of freedom, residual $n - p - 1$, total $n - 1$
- $MS_{\text{res}} = \hat{\sigma}^2 = 2.84$, so $RSE = 1.69$;
 $ESS = TSS - RSS = 5417.15 - 556.83$
- $R^2 = ESS/TSS = 4860.32/5417.15 = 0.897$; the F -statistic is the same one as on the previous slide

Model fit in multiple regression

model	R^2	RSE
TV	0.612	3.26
radio	0.332	4.27
TV + radio	0.897	1.68
TV + radio + newspaper	0.897	1.69

- $RSE = \sqrt{RSS/(n - p - 1)}$ and $R^2 = 1 - RSS/TSS$; here R^2 is also the squared correlation between $\mathbf{y}$ and $\hat{\mathbf{y}}$
- Adding radio to TV raises R^2 from 0.612 to 0.897 and cuts RSE from 3.26 to 1.68; adding newspaper changes nothing useful
- R^2 never decreases when a feature is added, useful or not, while RSE can increase (the divisor $n - p - 1$ shrinks); adjusted R^2 corrects this (next lecture)

Prediction Interval

Prediction and its uncertainty

- For a new observation $\mathbf{x}^* = (x_1^*, \dots, x_p^*)^\top$ that is not in the training data, with $\tilde{\mathbf{x}}^* = (1, x_1^*, \dots, x_p^*)^\top$, the prediction is

$$\hat{y}^* = \hat{\beta}_0 + \hat{\beta}_1 x_1^* + \dots + \hat{\beta}_p x_p^* = (\tilde{\mathbf{x}}^*)^\top \hat{\boldsymbol{\beta}}$$

- Uncertainty from estimating the coefficients (reducible): $\hat{f}$ differs from the true f ; the **confidence interval** covers the *average* response $f(\mathbf{x}^*)$ over many observations with these features
- Irreducible error: a single response is $y^* = f(\mathbf{x}^*) + \varepsilon^*$ even if f were known; the **prediction interval** covers an *individual* response and is always wider than the confidence interval
- Model bias: a linear model is an approximation of reality; the intervals assume it is correct (checked with residual plots in the next lecture)

Confidence and prediction intervals: formulas

- The variance of the prediction $\hat{y}^* = (\tilde{\mathbf{x}}^*)^\top \hat{\boldsymbol{\beta}}$ follows from $\text{Cov}(\hat{\boldsymbol{\beta}}) = \sigma^2 (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1}$:

$$\text{Var}(\hat{y}^*) = \sigma^2 (\tilde{\mathbf{x}}^*)^\top (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{x}}^*$$

- 95% **confidence interval** for $f(\mathbf{x}^*)$:

$$\hat{y}^* \pm t_{0.975, n-p-1} \hat{\sigma} \sqrt{(\tilde{\mathbf{x}}^*)^\top (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{x}}^*}$$

- 95% **prediction interval** for y^* :

$$\hat{y}^* \pm t_{0.975, n-p-1} \hat{\sigma} \sqrt{1 + (\tilde{\mathbf{x}}^*)^\top (\tilde{\mathbf{X}}^\top \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{x}}^*}; \text{ the extra 1 under the root is the variance of the new error } \varepsilon^*$$

Confidence and prediction intervals: simple regression

- For a new feature value x^* , the prediction is $\hat{y}^* = \hat{\beta}_0 + \hat{\beta}_1 x^*$, with variance

$$\text{Var}(\hat{y}^*) = \sigma^2 \left[\frac{1}{n} + \frac{(x^* - \bar{x})^2}{S_{xx}} \right], \quad S_{xx} = \sum_{i=1}^n (x^{(i)} - \bar{x})^2$$


- 95% **confidence interval** for $f(x^*) = \beta_0 + \beta_1 x^*$:

$$\hat{y}^* \pm t_{0.975, n-2} \hat{\sigma} \sqrt{\frac{1}{n} + \frac{(x^* - \bar{x})^2}{S_{xx}}}$$

- 95% **prediction interval** for y^* : $\hat{y}^* \pm t_{0.975, n-2} \hat{\sigma} \sqrt{1 + \frac{1}{n} + \frac{(x^* - \bar{x})^2}{S_{xx}}}$
- Both are narrowest at $x^* = \bar{x}$ and widen as x^* moves away from it; this is the special case $p = 1$ of the general formulas

Confidence and prediction intervals for Advertising

- Sales on TV at TV = 100:
 $\hat{y}^* = 11.79$, 95% CI [11.27, 12.30],
95% PI [5.34, 18.23]
- The confidence band is narrowest at the mean TV budget and widens towards the edges; the prediction band stays wide because ε^* dominates
- Adding radio = 20 (TV + radio model): $\hat{y}^* = 11.26$, CI [10.99, 11.53], PI [7.93, 14.58], narrower because RSE fell from 3.26 to 1.68



Qualitative Predictors

Qualitative features: dummy variables

- Credit data: average credit card balance (dollars) of $n = 400$ cardholders; do women and men differ?
- Encode gender with a **dummy variable**: $x^{(i)} = 1$ if person i is female, 0 if male; the first four cardholders are male, female, male, female:

$$\tilde{\mathbf{X}} = \begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 1 & 0 \\ 1 & 1 \\ \vdots & \vdots \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} 333 \\ 903 \\ 580 \\ 964 \\ \vdots \end{bmatrix}$$

- Model $y = \beta_0 + \beta_1 x + \varepsilon$: the mean balance is $\beta_0 + \beta_1$ for women and β_0 for men, so β_0 is the average for men (the **baseline**) and β_1 the average difference between women and men

Qualitative features with two levels: results

- The fitted values are the group means: $\hat{\beta}_0 = 509.80$ (SE 33.13) for men and $\hat{\beta}_0 + \hat{\beta}_1 = 509.80 + 19.73 = 529.54$ for women
- $\hat{\beta}_1 = 19.73$ (SE 46.05, $p = 0.669$): women carry about \$20 more, but the difference is not statistically significant
- The t -test of $H_0: \beta_1 = 0$ is the two-sample t -test for equal means with a pooled variance estimate
- Coding male as 1 and female as 0, or as ± 1 , changes what the coefficients mean but not the fitted values

Qualitative features with more than two levels

ethnicity	x_1 (Asian)	x_2 (Caucasian)
African American	0	0
Asian	1	0
Caucasian	0	1

- A feature with k levels needs $k - 1$ dummy variables; the level without one is the baseline; one 0/1/2 variable would impose an order
- $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \varepsilon$: β_0 is the baseline average, β_1 and β_2 are differences from the baseline
- Credit fit: $\hat{\beta}_0 = 531.00$, $\hat{\beta}_1 = -18.69$ ($p = 0.774$), $\hat{\beta}_2 = -12.50$ ($p = 0.826$)
- The individual p -values depend on the baseline; test the whole feature with the F -test: $F = 0.04$, $p = 0.957$

ANOVA for a qualitative feature

source	df	sum of squares	mean square	F	p -value
ethnicity	2	18,454	9,227	0.043	0.957
residual	397	84,321,458	212,397		
total	399	84,339,912			

- Regression on the $k - 1$ dummy variables reproduces **one-way ANOVA**: the regression sum of squares is the between-group sum of squares and the residual sum of squares is the within-group sum of squares (ANOVA means analysis of variance)
- The F -test of $H_0: \beta_1 = \beta_2 = 0$ is the test that all group means are equal; it does not depend on which level is the baseline
- Credit: the group means (531.0, 512.3, 518.5) differ far less than the person-to-person variation within groups, so $F = 0.043$ and $R^2 = 18454/84339912 \approx 0.0002$

Interaction

Interactions: relaxing the additive assumption

- The additive model assumes the effect of TV on sales is the same at every radio budget; spending on both together may work better than the sum of the parts (synergy)
- Add an **interaction term**, a new column of $\tilde{\mathbf{X}}$ with entries $x_1^{(i)} x_2^{(i)}$:
$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_1 x_2 + \varepsilon = \beta_0 + \beta_1 x_1 + (\beta_2 + \beta_3 x_1) x_2 + \varepsilon,$$
so the effect of x_2 depends on x_1

Interaction: example

- Advertising:
 $\widehat{\text{sales}} = 6.75 + 0.0191 \text{ TV} + 0.0289 \text{ radio} + 0.0011 \text{ TV} \times \text{radio}$; the interaction has $p < 0.0001$, R^2 rises from 0.897 to 0.968 and RSE falls from 1.68 to 0.94
- An extra \$1,000 on radio is associated with about 28.9 more units when $\text{TV} = 0$ and about 137.5 more units when $\text{TV} = \$100,000$
- **Hierarchical principle:** if an interaction is in the model, keep the main effects even when their p -values are large

Advertising with interaction: model summary


	coefficient	std. error	<i>t</i> - statistic	<i>p</i> - value
intercept	6.7502	0.2479	27.23	< 0.0001
TV	0.019101	0.001504	12.70	< 0.0001
radio	0.028860	0.008905	3.24	0.0014
TV × radio	0.001086	0.000052	20.73	< 0.0001

- The interaction is highly significant: the effect of radio on sales grows with the TV budget, $0.0289 + 0.0011 \text{ TV}$
- $R^2 = 0.968$, $\text{RSE} = 0.944$ on 196 degrees of freedom; $F = 1963.1$ on $(3, 196)$, $p < 0.0001$

Interaction with a qualitative feature

- Credit: balance on income and a student dummy; without the income $\times$ student interaction the two groups get parallel lines (same slope, different intercept)
- With it, the slope is $\hat{\beta}_1 = 6.22$ dollars per \$1,000 of income for non-students and $6.22 - 2.00 = 4.22$ for students
- But the interaction estimate -2.00 has SE 1.73 and $p = 0.249$: no strong evidence that the slopes differ

Interaction with a qualitative feature (cont.)



Interaction with a qualitative feature: model summary

	coefficient	std. error	<i>t</i> - statistic	<i>p</i> - value
intercept	200.62	33.70	5.95	< 0.0001
income	6.218	0.592	10.50	< 0.0001
student (yes)	476.68	104.35	4.57	< 0.0001
income × student	−1.999	1.731	−1.15	0.249

- Non-students (baseline): $\widehat{\text{balance}} = 200.62 + 6.22 \text{ income}$; students: $(200.62 + 476.68) + (6.22 - 2.00) \text{ income} = 677.30 + 4.22 \text{ income}$
- The interaction ($p = 0.249$) is not significant: no strong evidence that the income slopes differ between students and non-students
- $R^2 = 0.280$, $\text{RSE} = 391.6$ on 396 degrees of freedom; $F = 51.30$ on $(3, 396)$, $p < 0.0001$

Notation compared with the ISLP book

This slide index observations with a superscript and features with a subscript. The table maps this to the notation of ISLP (James et al.).

Quantity	This lecture	ISLP
Feature j of observation i	$x_j^{(i)}$	x_{ij}
Feature vector of observation i	$\mathbf{x}^{(i)}$	$\mathbf{x}_i$
Label of observation i	$y^{(i)}$	y_i
Residual of observation i	$e^{(i)}$	e_i
Feature matrix (no intercept column)	$\mathbf{X}$	$\mathbf{X}$
Design matrix (with column of ones)	$\tilde{\mathbf{X}}$	$\mathbf{X}$
New observation	$\mathbf{x}^*, y^*$	x_0, y_0

Questions?
